

METHODOLOGY & CLIENT BRIEF

The Synthetic-Marketer Audit

A standalone reference for prospects, clients, partners, and stakeholders who want to understand exactly what this audit measures, how it produces its findings, and why a business should commission one.

MAY 11, 2026

Contents

At a glance

Section 1 — What This Tool Does

The process, in plain language
Why this category of audit matters (market value)
What it is not

Section 2 — How This Tool Does What It Does

The methodology, end to end
Sources of strategic insight (the precedents the methodology draws from)
Observable evidence (the data the tool actually collects)
How the score is calculated
Quality assurance — the ten checks

Section 3 — Why This Tool Does What It Does the Way It Does

The benefits of consistent use
Why this is worth paying for

Section 4 — Anticipated Questions

Process and scope
Methodology and trust
What you receive
Action and value
Pricing and engagement

Closing — what to expect

A standalone reference for prospects, clients, partners, and stakeholders who want to understand exactly what this audit measures, how it produces its findings, and why a business should commission one.

This document is **not a per-engagement deliverable**. It is the durable explanation of the process — read it once, then trust the reports you receive on top of it.

At a glance

- **What you receive:** Nine documents per audit — one executive brief, one narrated walkthrough, one full scorecard, one competitor report, one audience/persona report, one 90-day implementation roadmap, one glossary, one dashboard PDF, and one bundled full-report PDF.
- **What it measures:** Six weighted categories — Content & Messaging (25%), Conversion Optimization (20%), SEO & Discoverability (20%), Competitive Positioning (15%), Brand & Trust (10%), Growth & Strategy (10%) — combined into a single composite score on a 100-point scale.
- **Turnaround:** 10–15 minutes of pipeline time per audit. Same-day delivery.
- **Inputs required:** A public URL. Nothing else. No analytics access. No backend credentials. No questionnaire.
- **Hard constraint:** Every audit is self-contained. Your report does not reference, name, or compare against any other client's audit. Ever.

Section 1 — What This Tool Does

The process, in plain language

Synthetic-Marketer takes any public website URL and produces a complete, defensible marketing audit — the kind of work that traditionally takes a senior strategist 6 to 8 hours of manual research, scoring, and writing. It does this in roughly fifteen minutes by sequencing automated discovery scripts, controlled AI analysis across five specialist analyst roles, and a mechanical quality-assurance pass that catches inconsistencies before any document is shipped.

The output is not a one-page screenshot, an SEO grade, or a list of generic best-practices. It is a complete, branded report set that a business owner can read in an evening, share with their team or vendor partners on Monday morning, and act on inside a week.

Why this category of audit matters (market value)

Every marketing decision a business makes — what to write on the homepage, what to spend on ads, which channels to expand into, which competitor to position against, where to invest the next quarter's content budget — is a bet. The quality of that bet depends on the quality of the underlying observation. Most businesses make those bets with very little structured observation. They go with instinct, with what worked last year, or with whatever the loudest voice in the room recommended.

A formal marketing audit replaces guessing with measurement. It surfaces:

1. **What is actually present** on the site versus what the operator believes is present.
2. **Where the conversion path is leaking** — friction points, missing trust signals, weak CTAs.
3. **How discoverable the brand is** in 2026's hybrid search landscape (Google, AI assistants, social).
4. **What competitors are doing better** and where the strategic opening sits.
5. **Who the buyer actually is** versus who the brand is currently speaking to.
6. **What to do next, in priority order, with effort and impact estimates.**

Historically, this type of audit was the entry point of every six-figure agency engagement. It was expensive because it was labor-intensive. The Synthetic-Marketer pipeline removes the labor cost without removing the rigor — making the same caliber of analysis available at a price point and turnaround time that fits a single quarter's discretionary marketing budget rather than a full-year retainer.

What it is not

To be clear about scope:

- It is **not a free SEO scanner**. Free scanners surface technical defects (broken links, missing meta tags, slow pages). This pipeline includes those checks, but it also produces strategic recommendations grounded in competitive context, audience definition, and message clarity — none of which a free scanner produces.
- It is **not a brand-strategy consulting engagement**. It diagnoses; it recommends; it ranks. It does not facilitate workshops, run user interviews, or produce new visual identity.
- It is **not implementation**. The roadmap tells you what to do and in what order. Doing the work is a separate engagement.

Section 2 — How This Tool Does What It Does

The methodology, end to end

The pipeline runs in **eight sequential phases**. Each phase has a defined input, a defined output, and an automated checkpoint before the next phase runs.

Phase	What happens	Why it matters
0	A date-stamped project bin is created	All artifacts live in one folder, named with the brand and audit date — easy to share, easy to archive
1	Discovery — the site is crawled, schema is validated, performance is measured via Google's PageSpeed Insights API, robots.txt is parsed for AI-crawler accessibility, content excerpts are pulled from every page	Every downstream conclusion is grounded in observed evidence, not assumption
2	Five parallel analyst roles review the discovery digest — Content & Messaging, Conversion Optimization, SEO & Discoverability, Competitive Positioning, and a combined Brand & Trust + Growth & Strategy reviewer	Specialization prevents a single analyst from over-weighting their preferred lens
3	A senior aggregator combines the five outputs into a unified scorecard with a composite score, Quick Wins, Strategic Recommendations, Long-Term Initiatives, and a Revenue Impact Summary	The composite is calculated mechanically — no analyst can hand-wave a final number
4	Two companion reports run in parallel — a deep Competitor Report (with comparison matrix) and an Audience Report (5–7 named personas with channel allocation and creative hooks)	Strategic recommendations only land when the operator can see who they're competing against and who they're talking to
5	Four supporting deliverables are written — Executive Brief, Walkthrough, Implementation Roadmap, and Glossary	Different stakeholders need different formats. A CEO reads the Brief; a designer reads the Walkthrough; an operations lead reads the Roadmap
6	Automated QA pass — ten mechanical checks for cross-audit references, score consistency across documents, broken internal links, missing files, placeholder content, and brand-name leakage	Nothing ships unless every check passes. This is the difference between a polished deliverable and a draft

Phase	What happens	Why it matters
7	A dashboard PDF is rendered with the score gauge, category breakdown bar chart, and findings table	The cover deliverable a client opens first
8	All seven markdown documents are rendered to per-document PDFs and bundled into a single FULL-REPORT.pdf with cover and table of contents	Clients can read on phone, laptop, or print — no markdown-reader required

Sources of strategic insight (the precedents the methodology draws from)

Synthetic-Marketer is not a novel invention; it is the codification of widely-validated frameworks into a repeatable pipeline. The methodology synthesizes:

- **Conversion Rate Optimization** — the BJ Fogg behavior model, MECLABS' conversion sequence heuristic, Joanna Wiebe's message hierarchy, the "5-second test" for hero clarity.
- **SEO and AI Search (GEO)** — Google's Quality Rater Guidelines (E-E-A-T), structured-data conformance per schema.org specifications, Core Web Vitals (LCP, INP, CLS) field thresholds, and the emerging norms for AI-crawler permissioning (`robots.txt` directives for GPTBot, ClaudeBot, PerplexityBot, Google-Extended, Applebot-Extended, and ~20 other identified AI agents).
- **Positioning** — April Dunford's obviously-awesome framework for category and differentiator definition, Al Ries / Jack Trout positioning canon for competitive set construction.
- **Audience definition** — the Jobs-to-be-Done taxonomy, Bryan Eisenberg's persona-mapping discipline, and standard Meta / Google / LinkedIn ad-targeting parameters as the practical output layer.
- **Brand & Trust** — the trust-stack hierarchy (named team, real photography, third-party review aggregation, social proof depth, certifications, press citations) used in conversion research since the early 2010s.
- **Growth & Strategy** — Reforge / Brian Balfour's growth-loop taxonomy, the standard channel mix audit, retention-signal analysis.

The pipeline does not invent new evaluation criteria. It applies the established ones consistently, every time, without skipping the ones that don't pattern-match to a particular analyst's habits.

Observable evidence (the data the tool actually collects)

Every score is grounded in mechanically-collected data. The discovery script captures, at minimum:

- **Site map data** — full sitemap.xml parse (including sitemap-of-sitemaps), robots.txt, CMS / platform fingerprint (Squarespace, WordPress, Webflow, Shopify, Next.js, etc.), all internal links from the homepage (capped at 25 pages for breadth without runaway).
- **Page-level technical signals** — title tags, meta descriptions, H1/H2/H3 hierarchy, image alt-text coverage, duplicate-meta detection, headings-missing flags.

- **Structured data** — every JSON-LD block extracted, all `@type` values catalogued, schema validated for known defects (NAME-as-address, generic LocalBusiness without subtype, malformed `openingHours`, PostalAddress missing `streetAddress`, Organization missing `sameAs`).
- **AI-search readiness** — `robots.txt` parsed for blocks against ~25 known AI crawlers; an AdsBot-blocked site (which damages Google Ads quality score) is flagged automatically.
- **Common-page probing** — ~25 standard slugs (`/blog`, `/pricing`, `/team`, `/faq`, `/events`, etc.) are probed for HTTP 404s, surfacing "missing pages" findings without speculation.
- **Business identity** — phone numbers, email addresses, postal addresses, hours of operation extracted from page text and structured data; address-consistency check.
- **Conversion surface** — every form on the site catalogued (action URL, method, field-by-field summary including type, required flag, placeholder).
- **External / social** — every linked social platform, third-party storefront (Shopify, Square, etc.), and review platform (Yelp, Trustpilot, BBB, G2) identified.
- **Content excerpts** — testimonials (extracted via class-keyword detection), blockquotes, hero copy (H1 + first substantial paragraph) per page. This last step closes a known blindness in metadata-only audits.
- **Performance metrics** — Google PageSpeed Insights / Lighthouse scores for Performance, Accessibility, Best Practices, and SEO across up to three conversion-relevant pages, plus lab metrics (LCP, CLS, TBT, TTI, FCP, Speed Index) and field metrics (CrUX p75 LCP, INP, CLS — real-world Chrome user data when available).

Every fetched page is saved to a raw HTML cache in the project bin, so any conclusion in the final report can be traced back to the source HTML on demand.

How the score is calculated

Each of the five analyst roles produces a category score on a 0–100 scale using a defined rubric:

Score band	Grade	Meaning
85+	A	Excellent — best-practice execution
70–84	B	Strong — competent execution with refinement opportunities
55–69	C	Adequate — functional but with clear gaps
41–54	D	Below average — material weaknesses
≤40	F	Critical — significant rework required

The composite score is the **weighted average** of the six category scores, rounded to the nearest integer:

- Content & Messaging × 25%

- Conversion Optimization × 20%
- SEO & Discoverability × 20%
- Competitive Positioning × 15%
- Brand & Trust × 10%
- Growth & Strategy × 10%

The weights reflect the practical reality that most websites lose more revenue to weak content and weak conversion paths than to any other category. SEO is weighted equal to conversion because organic discovery is the most cost-efficient acquisition channel when present, and the most expensive deficit when absent.

The composite score then appears identically across the Executive Brief, the Walkthrough, the Marketing Audit, the Implementation Roadmap, and the dashboard PDF. The QA pass refuses to ship the bin if these numbers ever disagree by even a single point.

Quality assurance — the ten checks

Before any audit is delivered, ten automated checks run:

1. **All required deliverables are present.** Nothing is missing.
2. **No cross-audit references.** No mention of prior audits, other clients, or comparative ranking against other engagements. Each audit reads as a self-contained document.
3. **No sibling-brand leakage.** If the operator has audited competitors or related brands, no other-bin brand names appear in this bin's deliverables.
4. **The composite score is consistent** across every document that reports it.
5. **All six category scores are present** in the Marketing Audit's breakdown table.
6. **Every internal markdown link resolves** to a real file in the bin.
7. **Phone numbers and identity claims are consistent** across the subject deliverables.
8. **No placeholder content** — no TODO, FIXME, lorem ipsum, [YOUR_*] template variables.
9. **The dashboard PDF exists.**
10. **All per-markdown PDFs and the bundled FULL-REPORT.pdf are rendered.**

A failed check exits the pipeline with a non-zero status. A failing audit is never shipped.

Section 3 — Why This Tool Does What It Does the Way It Does

The benefits of consistent use

A single audit gives a snapshot. The compounding value emerges from running audits **consistently** — quarterly on a focal brand, on every new client engagement, on every competitor before launching against them.

1. **A baseline that can be measured against.** Marketing investments are notoriously hard to evaluate. A scored audit gives you a calibrated starting point. When the score moves — up after a website redesign, down after a content team change, sideways after six months of "doing things" — you have a defensible, comparable measurement of whether the work is paying off.
2. **Compounding institutional knowledge.** Each audit is a permanent record. Six months from now, you can reread the Executive Brief and know exactly what the situation was when a decision was made. New team members can be onboarded in an hour by reading the audit set instead of re-deriving the context.
3. **Defensible recommendations.** Every recommendation in every report points back to specific, observable evidence — a quoted line of body copy, a missing schema field, a specific competitor's positioning, a measured Core Web Vitals number. When a stakeholder asks "why are we doing this?", the answer is in the report, with the source.
4. **Prioritization without politics.** The Implementation Roadmap ranks work by effort $\times$ impact. The next thing to do is the next thing on the list. This dissolves the recurring marketing-team argument about whose pet project gets attention next.
5. **Same caliber of analysis, every time.** The pipeline does not have good days and bad days. It does not skip the boring checks because it's been a long week. It does not over-weight whatever framework it learned most recently. The methodology is fixed; the consistency is mechanical.
6. **Speed of decision.** A 10–15-minute audit means a marketing decision that would have waited two weeks for a human deliverable can now be informed inside a single working session. That speed is the difference between acting on insight and remembering the insight after the moment has passed.

Why this is worth paying for

A business owner evaluating this service is implicitly asking: *what am I paying for that I couldn't get for free or cheaper?*

The honest answer:

- **You are paying for completeness.** Free scanners hand you a list of technical defects. They do not produce a competitor report, they do not define your audience, they do not write an implementation roadmap, and they do not connect the technical findings to revenue impact. The work of synthesizing all of those into one decision-ready document is the work you are commissioning.
- **You are paying for the absence of bias.** Hiring an SEO agency to audit your site produces an audit that recommends SEO work. Hiring a paid-ads agency produces an audit that recommends paid ads. The Synthetic-Marketer pipeline has no service to upsell — it scores all six categories with the same rigor and lets the data direct the priority.
- **You are paying for the methodology, not the keystrokes.** The value is not in the time it takes to generate the report; it is in the fact that the methodology was built, validated, and codified once, and is now applied to your site with the same discipline it was applied to every other site. You are buying the framework and the QA pass, not the runtime.
- **You are paying for what gets caught.** The QA layer is the part that distinguishes a deliverable from a draft. Ten mechanical checks catch the inconsistencies that erode trust in any document delivered by a person. No human-written audit gets ten passes of QA before being sent.
- **You are paying for an asset you can use.** The deliverable set is reusable: the personas inform every subsequent ad campaign, the competitor report informs every positioning conversation, the roadmap structures the next two quarters of marketing work. A single audit pays for itself the first time a recommendation in it changes a decision that would have cost more.

If a business owner is making any marketing decision worth more than a few thousand dollars in a quarter — which is essentially every business with a website — the cost of being wrong about that decision exceeds the cost of this audit by an order of magnitude. That is the value math.

Section 4 — Anticipated Questions

The 25 questions a thoughtful prospect or client is most likely to raise. Answered directly.

Process and scope

1. What kinds of websites does this work on? Any public website. The pipeline has been validated on independent retail (banking, freight finance, garden design), professional services, e-commerce, and small media properties. It works equally well on a 10-page Squarespace site and a 500-page WordPress site — the discovery layer auto-detects the platform and adapts.

2. Do you need access to my analytics, ad accounts, or backend? No. The audit runs entirely on the public-facing site. This is intentional. It keeps engagement friction at zero, preserves your data privacy, and produces a result that reflects what every visitor and every search engine actually sees.

3. How long does an audit take? 10–15 minutes of pipeline runtime. Same-day delivery. Most of that time is the parallel analyst phase and the performance-metric API calls. Wall-clock from request to delivery is typically under an hour.

4. Who actually does the analysis — a human or an AI? The analysis is performed by AI analyst roles operating against a mechanically-collected evidence digest, with all outputs validated by an automated QA layer. The methodology, prompts, scoring rubric, evidence-gathering checklist, and QA checks are human-designed and human-maintained. Think of the pipeline as a senior strategist's workflow encoded so it can run reliably and consistently — not as "AI replaces human judgment" but as "human judgment, captured once, applied every time."

5. Can you audit competitor sites for me? Yes. The same audit can be run on any public URL. Many engagements include audits of the subject brand plus 1–3 direct competitors so the strategic comparison is grounded in observed data rather than speculation.

Methodology and trust

6. How do you decide my score? Six categories are scored 0–100 by specialized analyst roles using a defined rubric. The composite is a fixed weighted average — Content 25%, Conversion 20%, SEO 20%, Competitive 15%, Brand & Trust 10%, Growth & Strategy 10%. Every score in every document is consistent because the QA layer refuses to ship the audit if it isn't.

7. Why these six categories and not others? These are the six categories that appear repeatedly in the established marketing literature — CRO, SEO, positioning, audience definition, brand trust, and growth strategy — and they are the categories whose deficits most reliably correlate with revenue underperformance. Other categories (PR, partnerships, product-led growth) are addressed inside the Strategic Recommendations when relevant but are not separately scored to keep the report focused.

8. How do I know this isn't generic AI-generated advice? Every recommendation in the report is grounded in evidence collected from your specific site — a quoted line of copy, a measured performance number, a specific schema defect, a named competitor. The reports include exact rewrites (current → suggested → why), specific page references, and measurable success criteria. Generic advice ("improve your SEO," "consider testimonials") is explicitly forbidden by the methodology and caught by the QA layer.

9. What happens if the audit is wrong about something? Disagreements are surfaced and resolved against the source evidence. Every fetched page is preserved in the audit's raw HTML cache, so any factual claim can be checked against the actual page content. If a recommendation doesn't fit your context (you've already tried it, or you have a strategic reason to do otherwise), the Implementation Roadmap is structured so individual items can be deprioritized without invalidating the rest.

10. How do you handle industries you don't know well? The methodology is industry-agnostic by design. The discovery layer collects observable evidence from the site itself, and the analyst roles reason about that evidence using established frameworks rather than industry insider knowledge. For highly specialized industries, the audit identifies the structural / strategic issues; industry-specific tactical depth is added through follow-up consultation if needed.

11. Is my data shared, stored, or used for AI training? No. The audit runs as a private operator-controlled process. Your project bin lives only in the operator's local environment unless you explicitly receive a copy. No data is shared with third parties beyond Google's PageSpeed Insights API (which receives only the public URL) and the Anthropic API (which receives the discovery digest for analyst reasoning, under Anthropic's commercial terms — not used for model training).

12. What gives this methodology authority? It is built on a specific stack of established frameworks: Google's E-E-A-T quality criteria, schema.org structured-data specifications, Core Web Vitals as defined by Chrome's CrUX, the BJ Fogg behavior model, April Dunford's positioning framework, Reforge's growth-loop taxonomy, and the standard CRO trust-stack hierarchy. The pipeline applies these consistently. The QA layer enforces consistency. The result is reproducible and defensible — two of the same audit on the same site on the same day will produce the same score.

What you receive

13. What exactly do I receive? Nine documents in a single zip or folder: 1. **Dashboard PDF** — score gauge, category bar chart, top findings (the cover deliverable). 2. **Executive Brief** — one-page TL;DR for stakeholders who will not read the full report. 3. **Walkthrough** — narrated, conversational tour of the audit findings. 4. **Marketing Audit** — full 6-category scorecard with detailed analysis and prioritized actions. 5. **Competitor Report** — 4–8 named competitor profiles plus comparison matrix. 6. **Audience Report** — 5–7 detailed personas with targeting parameters and creative hooks. 7. **Implementation Roadmap** — 90-day execution plan with effort × impact ranking. 8. **Glossary** — terminology reference calibrated to your specific report. 9. **Bundled Full-Report PDF** — all of the above in a single document with cover and table of contents.

14. Will my team understand the deliverables, or are these for marketers only? Each document is written for a specific reader. The Executive Brief is for non-marketing stakeholders (one page, plain language). The Walkthrough is for cross-functional teams (conversational, narrated). The Marketing Audit and Implementation Roadmap are for the marketing or operations lead who will direct the work. The Glossary defines every specialized term used in your specific report. No prior marketing background is required to act on the deliverables.

15. Can I share these reports with my agency or vendors? Yes — that is one of the intended uses. The reports are designed to be self-contained and shareable. Many clients use the Implementation Roadmap as a direct briefing document for their existing agency or implementation team.

16. Is there a presentation version for executives? The Dashboard PDF and the Executive Brief together function as the executive presentation layer. The Dashboard PDF in particular is built to be the cover deliverable — opening it gives a board-ready visual summary in under 30 seconds.

Action and value

17. After the audit, what do I do? Read the Executive Brief first (5 minutes). Then read the Implementation Roadmap (15 minutes). The Roadmap is structured in four phases — Week 1 quick fixes, Weeks 2–4 foundation, Month 2 content engine, Month 3 competitive positioning — with an effort × impact score on every item. Start at the top and work down.

18. Do you implement the fixes? Not as part of the audit itself. The audit is the diagnosis and the prescription. Implementation is a separate engagement, either with the operator's team, your existing agency, or a referred specialist. The Roadmap is structured so any competent implementation team — internal or external — can execute against it.

19. How much does it cost to fix what the audit identifies? The Implementation Roadmap rates every action by effort (1–3) and impact (1–3). A typical Quick-Win item (effort 1, impact 2–3) is a few hours of designer / copywriter time. A Strategic Recommendation (effort 2, impact 3) is a multi-week project. The Roadmap's effort scoring is calibrated so a marketing operator can scope the cost per item before deciding whether to commission the work.

20. What ROI should I expect? The audit includes a Revenue Impact Summary that quantifies the estimated upside from addressing the top findings. ROI math depends on the size of the business and the size of the gap — a small business with a critical conversion deficit may recover the audit cost on a single fix; a larger business may use the audit to prioritize a quarter of marketing investment more efficiently. The audit makes the relationship between "what to fix" and "what it's worth fixing" explicit, so the ROI calculation is yours to verify in your context.

21. How often should I re-audit? Quarterly is the recommended cadence for a business in active growth or active marketing investment. Annually is the minimum for any business that wants to track its own marketing health over time. Re-auditing also produces a measurable delta — you can see what moved, what didn't, and whether the work shipped between audits actually produced the intended outcome.

Pricing and engagement

22. What does this cost? Pricing is engagement-specific and depends on scope (single brand vs. brand + competitors), depth (Haiku model for fast/efficient audits, Opus for deepest reasoning), and any custom additions. The standard single-brand audit is positioned in the low-four-figures range — a fraction of the cost of a comparable agency engagement, with a fraction of the turnaround time. Specific pricing on request.

23. Why pay for this when I could ask ChatGPT myself? You could ask ChatGPT for marketing advice, and you would get a useful conversation. What you would not get: the structured discovery against your specific site, the mechanical extraction of 30+ specific data points (schema defects, performance numbers, alt-text coverage, duplicate metas, robots.txt AI-crawler status, etc.), the five specialized analyst roles applied to that data, the cross-document consistency enforced by an automated QA layer, the pre-rendered nine-document deliverable set, or the prioritized roadmap with effort/impact ranking. This is the difference between a conversation and a deliverable. You are paying for the deliverable.

24. What's the difference between this and a free SEO scanner? A free SEO scanner is one component of one of the six categories audited here (SEO & Discoverability, weighted 20%). It will identify some — not all — of what the SEO analyst role identifies. It will not produce competitive analysis, audience definition, content/message critique, conversion friction analysis, brand trust assessment, growth strategy review, or an implementation roadmap. A free SEO scanner solves a small slice of one quadrant of the question this audit answers.

25. Can I see a sample report before commissioning one? Yes. A redacted sample audit (anonymized brand, anonymized competitors, full deliverable structure) is available on request. The sample shows exactly what you receive — every document, every section, every level of specificity.

Closing — what to expect

If you commission an audit, the workflow is:

1. **You provide a URL.** That is the entire input.
2. **The pipeline runs.** 10–15 minutes of runtime.
3. **You receive the deliverable bundle.** Nine documents, all branded, all client-ready, in PDF and markdown.
4. **You read the Executive Brief and the Implementation Roadmap first.** Twenty minutes of reading time.
5. **You begin work** — internally, with your existing team, or with an implementation partner of your choice.

There is no questionnaire, no kickoff call required, no two-week wait, no scope-creep on the audit itself. The audit is the audit; what you do with it is a separate decision.

For pricing, sample reports, or to commission an audit, contact the operator directly.

This document describes the standing methodology of the Synthetic-Marketer audit pipeline. The methodology is reviewed and revised on a published cadence. Pipeline version, discovery script version, and QA script version are stamped on every deliverable so any audit can be matched to the exact methodology version that produced it.